

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

In [2]:

```
train3_f = 'train3_oddYr.txt'
test3_f = 'test3_oddYr.txt'
train5_f = 'train5_oddYr.txt'
test5_f = 'test5_oddYr.txt'

train3 = np.loadtxt(train3_f, dtype=int)
test3 = np.loadtxt(test3_f, dtype=int)
train5 = np.loadtxt(train5_f, dtype=int)
test5 = np.loadtxt(test5_f, dtype=int)
```

In [5]:

```
label_train3 = np.zeros(train3.shape[0])
label_train5 = np.ones(train5.shape[0])
```

In [6]:

```
# concatenate the label and the 64 * 1 vector together for random shuffling
train3 = np.concatenate((train3,label_train3[:,np.newaxis]),axis = 1)
train5 = np.concatenate((train5,label_train5[:,np.newaxis]),axis = 1)
```

In [12]:

```
#extract x and y from training samples
train = np.concatenate((train3,train5))
np.random.shuffle(train)
x_train = train[:, :-1]
y_train = train[:, -1]
print(x_train.shape,y_train.shape)

(1400, 64) (1400,)
```

In [13]:

```
#initialization of w
w = np.random.randn(64,1) / 100
```

In [14]:

```
# define sigmoid function
def sigmoid(w,x):
    z = np.dot(x,w)
    return(1/(1+np.exp(-z)))
```

For this problem, I will use the ML estimator to estimate the parameters w :

- For the loss function of this binary-classification problem, I use the cross entropy loss;
- The optimization algorithm I use is gradient descent;
- Maximum iterations of this algorithm is 1000, learning rate = $1 / 700$;

- Every 10 iterations, the loss and the corresponding error rate will be updated and recorded.

In [15]:

```
alpha = 0.2 / x_train.shape[0]
max_iter = 5000
loss_list = []
error_list = []
for i in range(max_iter):
    prob = sigmoid(w,x_train)
    temp = np.log(prob) * y_train[:,np.newaxis] + np.log(1-prob) * (1-y_train[:,np.newaxis])
    loss = -1 / x_train.shape[0] * np.sum(temp,axis = 0)
    if i%10 == 0:
        loss_list.append(loss)
        print('after ',str(i), ' iterations, the loss is ',str(loss[0]))
        prob_cur = sigmoid(w,x_train)
        y_cur = np.where(prob_cur > 0.5,1,0)
        error_rate = np.sum(np.absolute(y_train[:,np.newaxis] - y_cur),axis = 0) / x_train.shape[0]
        error_list.append(error_rate)
        print('error rate is ',str(error_rate[0]))
    temp1 = (y_train[:,np.newaxis] - prob) * x_train
    gradient = np.sum(temp1,axis = 0)
    w = w + alpha * gradient[:,np.newaxis]
loss_list = np.array(loss_list)
error_list = np.array(error_list)
```

```
after 0 iterations, the loss is 0.6925897569617911
error rate is 0.49642857142857144
after 10 iterations, the loss is 0.4482931798418664
error rate is 0.10857142857142857
after 20 iterations, the loss is 0.3591451399418551
error rate is 0.08928571428571429
after 30 iterations, the loss is 0.3123788785022366
error rate is 0.08071428571428571
after 40 iterations, the loss is 0.28301422271796883
error rate is 0.07785714285714286
after 50 iterations, the loss is 0.26259392999591896
error rate is 0.07357142857142857
after 60 iterations, the loss is 0.24742817219668017
error rate is 0.07142857142857142
after 70 iterations, the loss is 0.23563585170006418
error rate is 0.07071428571428572
after 80 iterations, the loss is 0.22615127023652543
error rate is 0.06714285714285714
after 90 iterations, the loss is 0.2183225812085373
error rate is 0.06607142857142857
```



In [29]:

```
x = np.linspace(0,5000,500)
```

As we can see from the plot, as the number of iterations increase, the loss generally decreases and converges. Also the error rate generally decreases along with the increase of iteration number.

In [30]:

```
fig=plt.figure(figsize=(30,10))
fig.add_subplot(1,2,1)
plt.plot(x,loss_list,'b')
plt.xlabel('number of iterations', fontdict={'family' : 'Times New Roman', 'size'
plt.ylabel('loss', fontdict={'family' : 'Times New Roman', 'size' : 12})
fig.add_subplot(1,2,2)
plt.plot(x,error_list,'g')
plt.xlabel('number of iterations', fontdict={'family' : 'Times New Roman', 'size'
plt.ylabel('error rate', fontdict={'family' : 'Times New Roman', 'size' : 12})
plt.show()
```



In [31]:

```
# optimized weight vector
for i in range(w.shape[0]):
    print('w%d = %.3f'%(i,w[i][0]),end = '\t')
    if i % 3 == 2:
        print()
```

```
w0 = -0.890    w1 = -1.392    w2 = -1.151
w3 = -1.099    w4 = -0.745    w5 = -0.773
w6 = 0.816     w7 = 1.700     w8 = 0.068
w9 = -0.093    w10 = 0.202    w11 = -0.069
w12 = -0.336   w13 = 0.685    w14 = -1.218
w15 = -1.276   w16 = 3.219    w17 = 1.363
w18 = 1.352    w19 = 0.220    w20 = 0.625
w21 = -1.914   w22 = -2.384   w23 = -2.429
w24 = 0.785    w25 = 0.407    w26 = 0.553
w27 = -0.268   w28 = -0.488   w29 = -2.156
w30 = 0.354    w31 = -0.029   w32 = 0.474
w33 = 1.041    w34 = 0.046    w35 = -0.318
w36 = -0.626   w37 = -0.205   w38 = -0.398
w39 = -0.322   w40 = 1.119    w41 = -0.188
w42 = -0.313   w43 = -0.062   w44 = 0.104
w45 = -0.807   w46 = 0.762    w47 = -1.410
w48 = 1.365    w49 = -0.596   w50 = 1.241
w51 = 0.554    w52 = 0.401    w53 = -0.291
w54 = 0.216    w55 = -1.132   w56 = 0.524
w57 = 0.277    w58 = 0.866    w59 = 1.668
w60 = 0.476    w61 = 0.628    w62 = 0.521
w63 = -0.461
```

In [32]:

```
# test data
label_test3 = np.zeros(test3.shape[0])
label_test5 = np.ones(test5.shape[0])
test3 = np.concatenate((test3,label_test3[:,np.newaxis]),axis = 1)
test5 = np.concatenate((test5,label_test5[:,np.newaxis]),axis = 1)
test = np.concatenate((test3,test5))
np.random.shuffle(test)
x_test = test[:, :-1]
y_test = test[:, -1]
print(x_test.shape,y_test.shape)
```

```
(800, 64) (800,)
```

In [33]:

```
# using the optimized weight vector to predict test data sample
prob_test = sigmoid(w,x_test)
pred_test = np.where(prob_test > 0.5,1,0)
error_rate = np.sum(np.absolute(y_test[:,np.newaxis] - pred_test),axis = 0) / x_test
print("testing error rate ",str(error_rate[0]))
```

```
testing error rate  0.06
```

In []: